

CSCE 581: Introduction to Trusted AI

Week 1 - Lectures 1 and 2: Introduction to AI, Trust and Real-World Applications

PROF. BIPLAV SRIVASTAVA, AI INSTITUTE


13TH AND 15TH JAN 2026

Carolinian Creed: “I will practice personal and academic integrity.”

Credits: Copyrights of all material reused acknowledged

Organization of Lectures 1, 2

- Introduction Section
 - Instructor introduction
- Main Section
 - Lecture 1: AI and Trust
 - AI: A quick introduction
 - (AI) Trust *
 - Expectations survey
 - Discussion: About the course
 - Related Courses: CSCE 580, 590s, 771
 - Course objectives and differentiation
 - Course logistics
 - Lecture 2: Data and Case Studies
 - Data analysis for traffic (South Carolina), Trust
 - Recommendations and Trust [Fairness and Teaming Recommendation]
- Concluding Section
 - About next lecture – Lecture 2
 - Ask me anything



Credit: Dilbert

Introduction Section



BIPLAV SRIVASTAVA


Research Snapshot (1989-2026)

Keywords: AI, Services, Sustainability

Current Research

Focus: **Theory** (Neuro-symbolic), **Usability** (Trust Rating, RCTs), **Smart Cities** (Energy, Water, Health)


The Space of AI Applications Explored



Details: <https://sites.google.com/site/biplavsrivastava/>
AI4Society Research Group: <https://ai4society.github.io/projects/>

Keywords: AI, Services, Sustainability
Papers: 250+ refereed; 7,200+ references
Patents: 77 (US issued); 4 sole inventions



Personal Programming Language Journey* (35+ years)



*Excluded: Javascript, Prolog, ...

Develop a Vibrant Research Culture Around AI

Students building chatbots
in Adv. Prog. Tech. class
in C++, Java and Python
(Elected Reps, Spring 2022;
Diseases, Spring 2023; Finance,
Spring 2024)



WebMD



AI/ Chatbots built for: governance (IJCAI 2016, AI Magazine 2024), **astronomy (AAAI 2018 best demo award)**, water (AAAI 2018), smart room (ICAPS 2018 demo runner up, IJCAI 2018), career planning (commercial product), **market intelligence (AAAI 2020 deployed AI award)**, dialogs for information retrieval (ICAPS 2021), fairness assessment (AAAI 2021), computer games (AAAI 2022), generalized planning (IJCAI 2024), **information spread in opinion networks (AAAI 2024 best demo award)**, transportation, set recommendation (teaming (AAAI 2024 deployed AI award), meals) and health.



System Image Credit:

Christine Steege, CSCE240(H), Spring 2023

Classes offered:

Trusted AI (CSCE 581)/ AI (CSCE 580) , Adv. Prog. Tech. (CSCE 240),
Comp. Proc. of Nat. Lang./NLP (CSCE 771)
Special Topics – Open Data, Planning, Chatbots

<https://ai4society.github.io/demos/>

TAI News

#1 NEWS – DOJ creates task force to challenge state AI regulations

Link - <https://www.cbsnews.com/news/doj-creates-task-force-to-challenge-state-ai-regulations/>

- **Context**

- US a backbencher in AI safety and governance; EU a leader
- Mixed culture in regulation – over in health, under in technology
- Federal govt restricting states from regulating
 - While not regulating at national level
- “A handful of states — including Colorado, California, Utah and Texas — have passed laws that set rules for AI companies, and several other states are considering regulations. Most AI-related bills proposed in state legislatures last year centered on protections from AI overreach, including legislation designed to restrict "deepfakes" and require companies to disclose if consumers are interacting with AI chatbots, [according to](#) the Brookings Institution”.

- **Implications**

- Thrusting technology (AI) without safety guarantee does not promote trust
- Risk mitigation is well understood (e.g., insurance for houses and cars); just not considered for AI


Main Section

AI: A Quick Introduction

Concept: AI

Example: Taking Care of a Baby

Individual's Extension



Expected behavior:


- Inform
 - Alert when crying
 - Alert when awake
 - Alert when idle
- Do
 - Raise temperature of room
 - Play music
 - ...

Conditions can be

- input and **reasoned** (e.g. **rule-based methods**) OR
- **learned** (from data)

Example: Taking Care of a Senior

Assisted Cognition




Expected behavior:

- Inform
 - Alert when idle
 - Alert when away from known locations
 - Alert when checkup/ medicines due
- Do
 - Send body parameters periodically
 - ...

Example: Taking Care of Oneself

Personal Digital Assistants




Expected behavior:


- Inform
 - When missing meetings
 - When missing social commitments
 - Reminding of priorities
 - ...
- Do
 - Make all cancellations / re-bookings when schedule changes
 - Find alternatives to current decisions and give choices (e.g., traffic)
 - ...

AI => Adaptive/ Intelligent Software System

- Business Logic Moves to Declarative Data (policy)
- Software is more resilient to changes in environment



Artificial Intelligence (AI) as an Agent



AI deals with perceiving the environment and taking actions towards short- and long term goals as the world changes over time.

Example 1: Courses for a Student

- Decision: Student deciding which courses to take for their program
- Data
 - **Public:** About courses
 - **Public:** About faculties
 - **Public:** About job opportunities
 - **Public:** About research opportunities and industry trends
 - **Private:** what the student wants to do
- Analysis
 - Courses offered in different semesters
 - Teachers offering courses – background, hardness of classes, ...

Trust

- Are the insights reliable?
- Do they cause short- or long-term harm?
- Will users adopt the insights?

Thought Exercise – (AI) Class and a Hypothetical AI-based Advisor

- **Good** decisions for students
 - Get good grades, marks
 - Learn
 - ...
- AI-may suggest
 - Give teacher rating
 - *But what about learning?*
- **Good** decisions for instructor
 - Get good rating
 - Finish course
 - Teach long-term skills
 - ...
- AI-may suggest
 - Give student grades
 - *But what about teaching?*

Trust

- Are the insights reliable?
- Do they cause short- or long-term harm?
- Will users adopt the insights?


Example 2: Health During a Pandemic

- Decision: Individual staying healthy during a pandemic like COVID19
- Data
 - **Public**: About disease, cases, deaths, variants
 - **Public**: About mitigation steps: e.g., mask wearing restrictions and practices, lockdowns, hospital conditions
 - **Private**: pre-existing health conditions
- Analysis
 - Regions with high and low cases
 - Whether to eat inside a restaurant?
 - How to make an urgent road trip ?
 - How to hold classes at a University?


Trust

- Are the insights reliable?
- Do they cause short- or long-term harm?
- Will users adopt the insights?

Advanced AI Techniques (**Analytics**) like Reasoning (**Symbolic**) & Machine Learning (**Neural**)
make use of data and models to provide insight to guide decisions



Analytics Landscape



Based on: Competing on Analytics, Davenport and Harris, 2007

History of Chatbots is the History of AI

Credit: https://en.wikipedia.org/wiki/Turing_test

1950 - Turing test

“which player – A or B – is a computer and which is a human.”

1964-66 – Eliza


computerized Rogerian psychotherapist

<https://en.wikipedia.org/wiki/ELIZA>, <http://www.manifestation.com/neurotoys/eliza.php3>

2011 – IBM Watson

question answering in a game setting


Today – Amazon Alexa, Google Echo, Apple Siri, ...



Credit: https://en.wikipedia.org/wiki/IBM_Watson

Trust, Trustworthiness and Trusted-AI

Values, Ethics, Law, Decisions (Actions) with Trust



Trust: *how can, humans or machines (AI) or in collaboration together, take decisions and actions following the same **values**, **ethical** principles and **laws** ?*

Assumption: *they should follow same values, principles and laws.*

Trust Scenario



Alan – wants to give money

Trust Scenario



Alan – wants to give money

- Could be first time or regular
- Wants to be effective and efficient


Decisions:

- Whom to give
- How much to give
- When to give

Trust Scenario



Alan – wants to give money



Trust Scenario



Alan – wants to give money

What decisions should be made by Alan?

Candidates

- Want money
 - May be more needy (or effective) than others
 - May be more efficient (less wasteful) than others in using it
- May change behavior after receiving donation
- May use money in different ways than promised

A Lesson in Trust

Weather alerts and Closing campus, Canceling classes

- Event order and response by actors // Choice 1 (Actual): Trustable ??
 - Alert1 -> Close campus -> Cancel class
 - Alert2 -> Unclose (Open) campus -> Uncancel (Normal) class
 - Alert3 -> Close campus -> Cancel class
 - ...
- Event order and response by actors // Choice 2: A more trustable way ??
 - Alert1 -> Close campus -> Online class (or recorded) OR CANCEL class
 - Alert2 -> Unclose (Open) campus -> No Change
 - Alert3 -> Close campus -> No Change
 - ...
- Which one would you have preferred, and WHY?

Why is Ethics Even an Issue with AI?

- When a technology works with humans and relates to inter-personal issues, the question of ethics comes into picture
- Examples: donations/ organs, medicine (opiods), food (genetically modified)

Discussion: what, if any issue,


- in recommending courses to students?
- in finding treatment for Covid?

Expectations Survey

- 9/15 = 60% response rate as of 6:45pm on Jan 12, 2026.
- AI background (1-5 scale):

I (instructor) want to know how good your background in AI is. Enter on a scale of 1-5 (5 highest). If this is your first course (e.g., not taken a course like Pattern Recognition before), please enter between 1-2.

9 responses



Main thing to learn

- How to use AI for future jobs
- Nothing in particular, I just want more general knowledge about AI
- the history and development of AI models
- Evaluate AI systems for transparency, fairness, and accountability
- I'm unsure, I just want a better general understanding of AI
- Understanding neural networks and how decisions are made.
- How AI is used in decision making
- How AI models validate their findings to be true or trustworthy.
- Creating less bias and safe AI for people.


AI problems of interest

- AI for entity behavior and pathfinding in video games
- No clue, I've done some outside AI courses, but this will be my first in a university setting, so I have no clue.
- Past AI models using propositional logic instead of probabilistic prediction (mentioned by a previous professor, sounded interesting)
- decision making process
- optimization
- Decision making under uncertainty and optimization
- Sports and fashion and arts
- Decision making under uncertainty would be interesting to learn about.
- Decision Making Uncertainty, Bias, and Security regarding LLMs.

Expectations Survey


GenAI/LLMs are quite proficient in generating content but they can also be problematic. How much would you like me (instructor) to allow you to use it? Enter on a scale of 1-5 (5 highest).

9 responses



GenAI/LLMs are quite proficient in generating content but they can also be problematic. How much would you like me (instructor) to use it for course material evaluation? Enter on a scale of 1-5 (5 highest).

9 responses



Trust Topics - (explanations, fairness (lack of bias), testing and rating; and in the domains of traffic, election, health and finance.)

- All of them
- Those sound cool
- AI's tendency to reinforce the status quo; how competitive AI is controlled by a few wealthy and powerful men
- Privacy
- unsure
- Not sure what exactly is all involved in trusted AI, but after looking up some other topics, **human oversight and data protection**.
- Some form of arts and fashion and paintings
- The domain of mathematics / problem solving
- Perhaps government services and how minorities are poorly represented and treated in the context of AI systems.

Expectations Survey

Any other comments?

- Just excited for the semester.
- Looking forward to great learning sessions
- For the last two questions. The 3 that I put for "how much would you like me to allow you to use it?", I want to have practice using GenAI/LLMs. For the final question, the 4 I put is meaning **I would like it to be used in lectures as a means to teach by example, so I get a better grasp on GenAI/LLMs.**
- Please help us make the class easy for us. I want to learn about Ai but please don't make it hard for us, **I speak for a lot of us when we say that we're really trying our best and already have a lot on our plates so please Professor, make it easy and entertaining for us** . Thank you in advance and I can't wait to enjoy this class
- I am very excited for the opportunity to take this class. I am especially interested to hear about not just the history and development of trustworthy AI systems, but **also the frontier research being done across the world to make these systems less bias and reliable.**

Course Logistics

Course Description – Spring 2026 (*)

CSCE 581 - Trusted Artificial Intelligence (3 Credits)

<https://cse.sc.edu/class/581>

AI Trust – responsible/ethical technology, fairness/ lack of bias, explanations (XAI), machine learning, reasoning, software testing, data quality and provenance, tools and projects.

Prerequisites: C or better in [CSCE 240](#) and [CSCE 350](#).

Prerequisite or Corequisite: D or better in [CSCE 330](#).

High Level Plan (Original)

CSCE 581 –

- Week 1: Introduction
- Week 2: Background: AI - Common Methods
- Week 3: The Trust Problem
- **Week 4: Machine Learning (Structured data) - Classification**
- **Week 5: Machine Learning (Structured data) - Classification – Trust Issues**
- **Week 6: Machine Learning (Structured data) – Classification – Mitigation Methods**
- **Week 7: Machine Learning (Structured data) – Classification – Explanation Methods**
- Week 8: Machine Learning (Text data) – Classification, **Large Language Models**
- **Week 9: Machine Learning (Text data) - Classification – Trust Issues**
- **Week 10: Machine Learning (Text data) – Classification – Mitigation Methods**
- **Week 11: Machine Learning (Text data) – Classification – Explanation Methods**
- Week 12: Emerging Standards and Laws
- Week 13: Project presentations
- Week 14: Project presentations, Conclusion

AI/ ML topics and with a
focus on fairness, explanation,
Data privacy, reliability

Reference: Intro AI Course Description

CSCE 580 - Artificial Intelligence (3 Credits)

Heuristic problem solving, theorem proving, and knowledge representation, including the use of appropriate programming languages and tools.

Prerequisites: [CSCE 350](#).

High Level Plan (Typical)

CSCE 580 – Introduction to AI – Topics in Recent Courses

- Topic 1: Introduction, aims
- Topic 2: Search, Heuristics
- Topic 3: Constraint Satisfaction Problems
- Topic 4: Decision making - Game trees
- Topic 5: Decision making - Decision networks
- Topic 6: Decision making – Markov Decision Processes, Hidden Markov models
- Topic 7: Learning – naïve Bayes, regression, Classification, clustering (unsupervised)
- Topic 8: Learning neural network, deep learning
- Topic 9: Decision making – Planning, Reinforcement Learning
- Topic 10: Robotics
- Topic 12: Representation, Ontology
- Topic 12: Tools

Fall 2024

Classical AI topics and a
focus on implementation

Reference: AI Learning Objectives

Understand the breadth of AI techniques, be empowered to solve real-world challenges

- L1: Appreciate and work with diversity of data— text, speech and visual; focus of course will, be structured data (e.g., tables) and text (NLP; English)
- L2: Learn techniques to derive insights from data spanning reasoning (e.g., symbolic) and learning (e.g., neural) in a decision-making setup
- L3: Learn methods to represent and organize insights
- L4: Make insights usable with people in a collaborative setting (“chatbots”)
- L5: Understand issues related to usage of AI methods/ tools with people.
- L6: Gain experience by build a real-work AI

Adapt Based on Class Interest?

- CSCE 581: AI/ ML topics and with a focus on fairness, explanation, Data privacy, reliability
- CSCE 580: Classical AI topics and a focus on implementation
- Need to adapt?
 - AI/ ML topics with a focus on generative AI, fairness, explanation, adversarial attacks; building chatbots

Administrative Information – CSCE 581

Walk through of
Github:

<https://github.com/biplav-s/course-tai-s26/>

Website:

<https://sites.google.com/site/biplavsrivastava/teaching/ai-csce-581-spring-2026-trusted-ai>

Quick Info - When and Where

- Tuesday/Thursday 11:40 am – 12:55 pm
- In person at **Swearingen Engr Ctr| Room 2A15**. Recordings will be available on Blackboard whenever feasible.

Catalog Information

- [Trusted AI - CSCE 581 001](#)
- CRN: 53067
- Duration: 01/12/2026 - 05/06/2026

Instructor Information

- Instructor: Biplav Srivastava
- E-mail: [biplav.s AT sc.edu](mailto:biplav.s@sc.edu)
- Office Hours: 10:15-11:15am (Tu, Th); other times by appointment
- [GitHub](#) for slides, sample code.

Course Material

- Artificial Intelligence: A Modern Approach (Fourth edition, 2020), Stuart Russell and Peter Norvig,
 - <http://aima.cs.berkeley.edu/>,
ISBN-13: 978-0134610993
- Trustworthy Machine Learning, by Kush R. Varshney,
<http://www.trustworthymachinelearning.com/>, 2022

Open Datasets

- data.gov from ANY COUNTRY
 - Portal: <https://dataportals.org/>
 - US: <https://www.data.gov/> or any US state
 - India: <https://data.gov.in>
- Text of legislations - LegiScan, <https://legiscan.com/>
- Kaggle datasets: <https://www.kaggle.com/datasets>
- Google datasets search:
<https://datasetsearch.research.google.com/>


- AI Fairness
 - Trisha Mahoney, Kush R. Varshney, and Michael Hind, Available at:
<https://krvarshney.github.io/pubs/MahoneyVH2020.pdf>
 - In AI We Trust: Ethics, Artificial Intelligence, and Reliability, Mark Ryan.
Available at: <https://link.springer.com/article/10.1007/s11948-020-00228-y>
- Python for Data Analysis
 - Latest: Python for Data Analysis Book, by Wes McKinney, 2nd Edition. On Amazon at: <https://www.amazon.com/gp/product/1491957662/>, ISBN-13: 978-1491957660, ISBN-10: 1491957662
 - Book Data and Code Notebooks: <https://github.com/wesm/pydata-book>
 - 1st edition (free download): <https://bedford-computing.co.uk/learning/wp-content/uploads/2015/10/Python-for-Data-Analysis.pdf>

Student Assessment

A = [920-1000]
B+ = [870-919]
B = [820-869]
C+ = [770-819]
C = [720-769]
D+ = [670-719]
D = [600-669]
F = [0-599]

Tests	Undergrad	Grad
Course Project – report, in-class presentation	600	600
Quiz – best of 2 from 3	200	200
Final Exam	200	100
Additional Final Exam – Paper summary, in-class presentation		100
Total	1000 points	1000 points

AI for the Real World



Credit: Dilbert – June 30, 2016

Lecture 2: Data

- Data analysis for traffic (South Carolina), Trust – <https://ai4society.github.io/projects/traffic-page/index.html>
- Recommendations and Trust [Fairness and Teaming Recommendation] – https://ai4society.github.io/projects/group_rec/index.html


Lecture 2 – Open Data

Lecture 2 - Outline


- Revisiting Intelligence
- (Open) Data
- Large Language Models and Assessing their Outputs
- Project Discussion

Intelligence: Different Shapes, Different Standards

- Exists regardless of body form, brain in nature
- May be task specific
- May involve a range of capabilities



Human tracheal skin cells self-assemble into multi-cellular, moving organoids called anthrobots. These images show anthrobots with cilia on their surface (*yellow*) distributed in different patterns. Surface patterns of cilia are correlated with different movement patterns: circular, wiggling, long curves or straight lines. Gizem Gumuskaya et al., "[Motile Living Biobots Self-Construct from Adult Human Somatic Progenitor Seed Cells](#)," *Advanced Science*, November 30, 2023



AI: Many Paths and Techniques

By goal pursued by an AI approach

- Building intelligent systems
- Understanding human brain / thought process – Neuroscience/ Cognitive Science
- Mimicking human behavior
- As advanced analytics (descriptive, predictive, prescriptive)

Legend: Areas we work in

By disciplines

- Representation
- Reasoning
- Learning
- Interaction (with people, other agents, ...): Chatbots
- Trust

AI4Society: *Innovations with
common man in mind!*



<https://ai4society.github.io/demos/>

Types of Data

- By media: Text, Sound (speech), Visual (image, video), Multi (modal, media)
- By structure: unstructured, semi-structured, structured
- By features: time-series, labeled/ unlabeled, spatio-temporal,





Image credit:


<http://www.trustworthymachinelearning.com/trustworthymachinelearning-04.htm>

Open Data

- Open data is the notion that data should not be hidden, but made available to everyone to **reuse**. **The idea is not new.**
- Scientific publications follow this: “standing on the shoulders of giants”
- Data quality and open publishing process is critical




USA



India

Does Opening Data Make It Reusable? No

Illustration of Levels



4

Temperature forecast for Galway, Ireland	
Day	Lowest Temperature (°C)
Saturday, 13 November 2010	2
Sunday, 14 November 2010	4
Monday, 15 November 2010	7

1

	A	B
1	Temperature forecast for Galway, Ireland	
2		
3	Day	Lowest Temperature (°C)
4	Saturday, 13 November 2010	2
5	Sunday, 14 November 2010	4
6	Monday, 15 November 2010	7

2

Temperature forecast for Galway, Ireland

Day	Lowest Temperature (°C)
Saturday, 13 November 2010	2
Sunday, 14 November 2010	4
Monday, 15 November 2010	7

5



```
gtd-3.csv - WordPad
File Edit View Insert Format Help

"Temperature forecast for Galway, Ireland",
"Day","Lowest Temperature (C)"
"Saturday, 13 November 2010",2
"Sunday, 14 November 2010",4
"Monday, 15 November 2010",7
```

3

Source: <http://5stardata.info/>

About 600 Data Catalogs of Public Data



Source: IBM Institute for Business Value.

As on 17 Aug 2022

Guideline: Human Impact of AI

- We study technology (AI) but it works with data
- Data, when from people or about people, can have issues like bias
 - **Example:** data reveals a view which is influenced by data collection practices
 - **Difference:** **World as it is**, world according to data and **world as it should be**
- The course and instructor believes in
 - Not promoting bias of any kind
 - Respecting everyone regardless of background

AI Ethics

Why is Ethics Even an Issue?

- When a technology works with humans and relates to inter-personal issues, the question of ethics comes into picture
- Examples: medicine (opioids), food (genetically modified)

Discussion: what, if any issue,

- in recommending courses to students?
- in finding treatment for Covid?

What is Specific to AI?

- AI needs **data**
 - Data privacy and governance
- AI is often a **black box**
 - Explainability and transparency
- AI can make **decisions/recommendations**
 - Fairness and value alignment
- AI is based on statistics and has always a small percentage of **error**
 - Who is accountable if mistakes happen?
- AI can infer our preferences and **manipulate** them
 - Human and moral agency
- AI is very **pervasive and dynamic**
 - Larger negative impacts for tech misuse
 - Fast transformation of jobs and society

Credits:

Tutorial on [Trusting AI by Testing and Rating Third Party Offerings at IJCAI 2020](#), Biplav Srivastava, Francesca Rossi, Jan 2021

Main AI Ethics Issues



DATA GOVERNANCE
AND PRIVACY



FAIRNESS AND
INCLUSION



HUMAN AND
MORAL AGENCY



VALUE ALIGNMENT



ACCOUNTABILITY



TRANSPARENCY AND
EXPLAINABILITY



TECHNOLOGY
MISUSE

Credits:

Tutorial on [Trusting AI by Testing and Rating Third Party Offerings at IJCAI 2020](#), Biplav Srivastava, Francesca Rossi, Jan 2021

Collaborative Assistants

- Conversation agents and interfaces (chatbots) are getting easy to build and deploy
 - Can be text-based or speech-based
 - Usually multi-modal (i.e, involving text, speech, vision, document, maps)
- Current chatbots typically interact with a single user at a time and conduct
 - Informal conversation, or
 - Task-oriented activities like answer a user's questions or provide recommendations

Demonstrations

- *Eliza*, <http://www.manifestation.com/neurotoys/eliza.php3>
- *Mitsuku*, <https://www.pandorabots.com/mitsuku/>
- ChatGPT, <https://openai.com/blog/chatgpt>

Exercise: Session with ChatGPT

- Ask questions about Water usage
 - Experience
- Ask questions about Finance
 - Experience
- Hint:
 - Demand / supply questions: “can I drink water of Lake Murray”?, “will US have money to pay debt next year”
 - Decision questions: “which water should I choose between a bottled one and tap”?
 - Factoid questions: “is pH of 7 good for drinking water?”

Exercise: Your Resumes

- What does a search (Google search) tell about you?
- What does a LLM/ ChatGPT tell about you?
- Task:
 - Put your resume at: <TBD>
- Course task: We will analyze them as part of AI/ data science activity in a later class

Exercise: Solving Games with AI

- Popular way to learn AI is via games
 - <https://github.com/biplav-s/course-ai-tai-f23/blob/main/sample-code/Class1-games.md>

Exercise: Comparing Outputs of GenAI


Consider llm_responses:

- "Google": "Title: Jimmy Kimmel Reacts to Donald Trump Winning the Presidential ... Snippet: Nov 6, 2024 ...",
- "Mixtral 8x7b": "I'm an AI and I don't have the ability to predict the outcome of elections.",
- "SafeChat": "Sorry, I am designed not to answer such a question.", }

Exercise: Comparing Outputs of GenAI


Consider llm_responses:

- "Google": "Title: Jimmy Kimmel Reacts to Donald Trump Winning the Presidential ... Snippet: Nov 6, 2024 ...",
- "Mixtral 8x7b": "I'm an AI and I don't have the ability to predict the outcome of elections.",
- "SafeChat": "Sorry, I am designed not to answer such a question.", }



Exercise: Comparing Outputs of GenAI

- General Approach (for any modality)
- Can be used for model pipelines as well



Exercise: Comparing Outputs of GenAI

- **(Python) Tool:** <https://pypi.org/project/GAICo/>
- **Repository:** github.com/ai4society/GenAIResultsComparator
- **Documentation:** ai4society.github.io/projects/GenAIResultsComparator

- **Code sample:**
<https://github.com/biplav-s/course-ai-f25/tree/main/sample-code/class2-gaico-usage>

Project Discussion

- Class inputs on individual projects
- Group project (?)
 - Kaggle competition - Google / Kaggle competition around MedGemma launched today - <https://www.kaggle.com/competitions/med-gemma-impact-challenge>. Models: (<https://huggingface.co/collections/google/health-ai-developer-foundations-hai-def>)

Concluding Section

Week 1: Concluding Comments

- We did a quick overview of
 - AI
 - Trust issues
- Course will focus on
 - Understanding trust issues and ongoing ways to make AI reliable, practical ways to convey trustworthy results to users.
 - Student evaluation will be by via project, paper and quizzes
- Exciting techniques to learn to impact the world around us

About Next Week – Week 2 (L3, L4)

High Level Plan (Original)

CSCE 581 –

- Week 1: Introduction
- Week 2: Background: AI - Common Methods
- Week 3: The Trust Problem
- Week 4: Machine Learning (Structured data) - Classification
- Week 5: Machine Learning (Structured data) - Classification – Trust Issues
- Week 6: Machine Learning (Structured data) – Classification – Mitigation Methods
- Week 7: Machine Learning (Structured data) – Classification – Explanation Methods
- Week 8: Machine Learning (Text data) – Classification, **Large Language Models**
- Week 9: Machine Learning (Text data) - Classification – Trust Issues
- Week 10: Machine Learning (Text data) – Classification – Mitigation Methods
- Week 11: Machine Learning (Text data) – Classification – Explanation Methods
- Week 12: Emerging Standards and Laws
- Week 13: Project presentations
- Week 14: Project presentations, Conclusion

AI/ ML topics and with a
focus on fairness, explanation,
Data privacy, reliability

Lecture 3, 4:

- Trust Motivation, Review Scope, Data
- Data, Knowledge Graph
- Project Timelines

Class #	Date	Description	Comments
1	Jan 13 (Tu)	Introduction, Trusted AI	W1
2	Jan 15 (Th)	Case Studies: Data Analysis for AI, Analysis for Trust [Traffic], Recommendations and Trust [Fairness and ULTRA]	
3	Jan 20 (Tu)	Review: Trusted Decisions, Expectations, Course Scope; Data	W2
4	Jan 22 (Th)	AI: Data Prep, Knowledge Graph	
5	Jan 27 (Tu)	Common AI methods: ML Landscape	W3
6	Jan 29 (Th)	AI - Structured: Analysis – Supervised ML	